

Highlights

Two-stage MILP scheduling and structure-aware learning-enhanced branch-and-cut for dual-source mixed-flow rotating-chamber cluster tools

Anonymous author(s)

- A two-stage MILP unifies dual-source mixed rotating-chamber scheduling.
- Makespan-preserving refinement sharply improves waits and chamber cadence.
- SPBS cuts mean PDI by 26.19% and raises incumbent availability to 100%.
- Variable-role reconstruction makes beam cut decoding structure-aware.
- RDMCT-HBS has the lowest observed mean PDI and time over 2700 nominal runs.

Two-stage MILP scheduling and structure-aware learning-enhanced branch-and-cut for dual-source mixed-flow rotating-chamber cluster tools

Anonymous author(s)

Abstract

Dual-source mixed-flow rotating-chamber cluster tools coordinate product and reusable vacuum-zone dummy wafers across shared chambers, load locks, and robots. Scheduling couples recipe grouping, rotation, cleaning, material availability, and residence limits; a throughput-optimal sequence can still exhibit poor waits and cadence. We formulate the first finite-horizon two-stage MILP integrating these decisions in mixed $4 \times 1/2 \times 2$ operation. Stage one minimizes end-to-end makespan; stage two retains the discrete schedule and makespan within numerical tolerance while refining physical route waits and chamber cadence. To solve it, we develop rotating dual-source mixed-flow cluster-tool hierarchical beam search (RDMCT-HBS), combining a structure-preserving binary-skeleton warm start (SPBS), variable-role reconstruction, A3C-inspired parallel policy-gradient learning, bounded beam decoding, and structural cut completion inside SCIP. Four exact cases verify the formulation. In 2700 nominal runs, RDMCT-HBS achieves the lowest mean primal-dual integral and solution time among seven methods, improving mean PDI by 1.41% over adapted HEM and 1.97% over SCIP. A separate 960-run comparison shows that SPBS reduces mean PDI by 26.19% and raises incumbent availability from 51.67% to 100%. In 1500 controlled runs, the second stage reduces physical route-total wait, route-maximum wait, cadence variation, and cadence deviation by 51.91%, 68.89%, 39.70%, and 94.05%, respectively. Across 729 OOD runs on nine larger configurations, all methods return an independently validated incumbent. The results establish both the operational value of the two-stage model and the nominal-range computational advantage of structure-aware exact search.

Keywords: rotating-chamber cluster tool, semiconductor manufacturing, mixed-integer linear programming, learning-enhanced branch-and-cut,

1. Introduction

Semiconductor equipment is capital intensive, and its effective capacity is often determined by scheduling rather than nominal processing speed. Cluster tools illustrate this dependence particularly clearly: chambers, load locks, and robotic handlers form a tightly coupled system in which a locally sensible decision can delay several downstream operations. Better scheduling therefore offers a direct route to higher asset utilization without additional hardware (Chien and Lan, 2021; Hong et al., 2026).

This paper studies a rotating-chamber cluster tool used in a plasma-enhanced chemical vapor deposition setting. The equipment receives material from two sources. Product wafers enter from atmospheric load ports, while reusable vacuum-zone dummy wafers enter from a dedicated storage area. The flows merge inside shared rotary chambers and compete for vacuum transport. Two recipe families coexist: 4×1 operation processes two chamber sides as a rotary batch, whereas 2×2 operation links consecutive product pairs through overlapping exposures. Cleaning creates additional dummy-only operations and splits processing into epochs. A schedule must consequently decide not only when to process, but also how to group products, which chamber and side to use, where to place cleaning boundaries, which physical dummy wafer to allocate, and how to order every conflicting robot and module action.

Existing cluster-tool research provides important models for cyclic operation, multiple wafer types, residency, cleaning, and multi-space modules (Qiao et al., 2022; Li et al., 2022; Wang et al., 2023a; Li and Yang, 2025; Yang et al., 2025; Gu et al., 2024). Those models do not jointly represent the finite-horizon dual-source flow, mixed rotary recipes, reusable dummy inventory, and shared transport considered here. The difference is structural. A prescribed robot sequence can be evaluated by an LP, but the present equipment requires the sequence, material assignments, and event times to be determined together.

Learning-based scheduling offers powerful search guidance, but most studies act directly on dispatching or process control (Liu et al., 2022; Zhang et al., 2024; Kim and Lee, 2025; Kim and Kim, 2025). The closely related CIE study of Hong et al. (2026) uses DDPG and discrete-event simulation to

35 adjust continuous processing times in a PVD cluster tool. Our problem is different: physical processing times are given, while the method must construct a feasible discrete schedule and retain valid lower bounds and optimality proofs. This motivates learning inside an exact solver rather than replacing the solver.

40 The paper resolves two linked gaps. The modeling gap is addressed by a two-stage MILP that separates throughput from schedule quality without separating the physical system. The computational gap is addressed by embedding equipment structure in both primal initialization and cut selection. The resulting research logic is direct: the dual-source rotary physics defines
45 the MILP; the MILP exposes its search bottlenecks; and those bottlenecks define the information used by the learning-enhanced branch-and-cut method.

The contributions are threefold.

1. We formulate a finite-horizon MILP for dual-source mixed-flow rotating-chamber scheduling. It jointly decides $4 \times 1/2 \times 2$ grouping, chamber and
50 side allocation, rotary events, robot transfers, load-lock states, cleaning epochs, residency, and reusable vacuum-zone dummy-wafer allocation.
2. We introduce a makespan-preserving second stage. Unlike makespan-only scheduling, it explicitly reshapes physical product waits and chamber cadence within the discrete schedule retained from stage one.
- 55 3. We develop RDMCT-HBS, a structure-aware enhancement of HEM that couples a feasible binary-skeleton warm start, variable-role cut features, A3C-inspired parallel learning, bounded beam decoding, and structural set completion while leaving feasibility and proof to SCIP.

The evidence design links the two contributions to complementary tests.
60 Exact verification establishes formulation fidelity; the nominal comparison measures algorithmic performance; controlled contrasts identify the warm-start and second-stage effects; and factorial and shifted-instance studies reveal the manufacturing conditions that govern computational difficulty and scalability.

65 **2. Research background and positioning**

2.1. Cluster-tool scheduling models

Cluster-tool models have progressed from steady-state robot cycles to richer representations of product heterogeneity and auxiliary constraints. In-

egrated patterns and workload balance have been investigated for dual-arm
 70 tools and multiple wafer types (Kim et al., 2021; Ko et al., 2021; Zhu et al.,
 2022). Recent MIP and analytical studies coordinate release order, robot
 tasks, chamber assignment, and residency for concurrent or multi-type pro-
 cessing (Li and Yang, 2025; Yang et al., 2025). Cleaning, purge, and post-
 processing constraints further demonstrate that temporal feasibility cannot
 75 be inferred from chamber capacity alone (Qiao et al., 2022; Lee et al., 2023;
 Yuan et al., 2023; Zhu et al., 2023; Pan et al., 2024).

Rotating and multi-space modules introduce side-dependent loading and
 shared chamber motion. Gu et al. (2024) analyse multi-finger robotic tools
 with multi-space process modules, providing the closest physical foundation
 80 for the chamber mechanism studied here. Their predetermined cyclic se-
 quences do not cover finite mixed batches or a second reusable material
 stream. The present problem extends the decision boundary from timing
 a given sequence to simultaneously constructing material groups, chamber-
 side assignments, cleaning epochs, resource orders, and event times.

85 2.2. Learning-enhanced scheduling and exact optimization

Deep reinforcement learning has been combined with genetic, adaptive,
 and problem-specific search in semiconductor scheduling (Chien and Lan,
 2021; Liu et al., 2022; Zhang et al., 2024; Li and Chang, 2022). These
 methods show that learned policies are most useful when coupled with do-
 90 main structure. Exact optimization adds a complementary advantage: ex-
 plicit constraints, primal and dual bounds, and verifiable feasibility. Recent
 learning-enhanced MIP research guides branching, primal search, and cut
 selection without surrendering those properties (Bengio et al., 2021; Khalil
 et al., 2022; Scavuzzo et al., 2024).

95 Cut selection is especially relevant because valid inequalities differ in both
 immediate efficacy and the decisions they constrain. Reinforcement learning,
 supervised ranking, look-ahead imitation, and adaptive scoring have all im-
 proved cut management (Tang et al., 2020; Huang et al., 2022; Paulus et al.,
 2022; Turner et al., 2023). HEM represents selection as a hierarchical count-
 and-sequence decision (Wang et al., 2023b, 2024). However, generic cut fea-
 100 tures cannot distinguish an inequality acting mainly on chamber assignment
 from one acting on continuous event timing; greedy decoding also discards al-
 ternative partial sequences. These limitations are particularly consequential
 in the proposed MILP because discrete equipment choices and continuous
 105 timing are strongly coupled.

Table 1 summarizes the research position. The novelty is not the isolated use of MILP, reinforcement learning, or beam search. It is the integration of a new dual-source rotary scheduling model with a solver method whose primal and dual guidance are both derived from that model’s structure.

Table 1: Research positioning relative to representative cluster-tool studies.

Study	Equipment setting	Decision scope	Solution approach
Qiao et al. (2022)	Cleaning-constrained tool	Residency and cleaning	Binary IP
Gu et al. (2024)	Multi-space rotary modules	Timing of preset sequences	LP
Yang et al. (2025)	Two wafer types	Release, robot tasks, residency	MIP
Kim and Lee (2025)	Concurrent cluster tool	Release and robot sequence	DRL + search
Hong et al. (2026)	PVD cluster tool	Continuous process-time control	DES + DDPG
This study	Dual-source mixed rotary tool	Grouping, material allocation, resources, cleaning, and event times	Two-stage MILP + learned B&C

110 3. Industrial problem and two-stage MILP

3.1. System logic and material interaction

Figure 1 shows the system at the level needed to understand the optimization model. Product wafers traverse the atmospheric domain, upper load lock, vacuum robot, rotary chamber, lower load lock, and return path. 115 Dummy wafers begin and end in the vacuum domain. The two streams share the vacuum transfer robot (VTR) and chambers but have different release and return requirements; the atmospheric transfer robot (ATR) handles the external product-wafer route.

The recipe determines how the flows interact. In 4×1 mode, the chamber 120 loads a front side, rotates, loads a back side, processes both sides, and unloads them in reverse order. A product-deficient final side is completed




Figure 1: Dual-source material flows and shared resources. Solid arrows denote the product-wafer route; dashed arrows denote reusable vacuum-zone dummy wafers.

by dummy wafers. In 2×2 mode, product pairs form a chain: adjacent exposures share an intermediate pair, while leading and trailing positions require dummy pairs. Cleaning interrupts the chain only at a legal exposure boundary and uses a dummy-only chamber cycle. A physical dummy wafer cannot be assigned to overlapping chamber occupancies.

This interaction creates three decision layers. The *material layer* chooses recipe groups, chambers, sides, cleaning epochs, and physical dummy wafers. The *resource layer* orders all operations competing for robots, load locks, and chambers. The *timing layer* assigns event times and enforces residence limits. These layers must be optimized jointly because a material choice changes both the resource conflicts and the feasible timing network.

3.2. Decision structure and principal constraints

Let $\mathcal{W}$ denote product wafers, $\mathcal{P}$ the product groups, $\mathcal{C}$ the rotary chambers, $\mathcal{K}$ the candidate recipe positions, $\mathcal{O}$ the operations, and $\mathcal{R}$ the unit-capacity resources. $\mathcal{K}(p)$ is the set of positions admissible for group p , $\mathcal{A}$ is the technological-arc set, and $\mathcal{O}_r$ contains the operations that may use resource r . For chamber c , $\mathcal{J}_c^{\text{DW}}$ is the set of dummy-wafer demands and $\mathcal{T}_c$ is its physical dummy inventory. Binary variables x assign product groups to recipe positions, u and b activate positions and operations, a assign process cycles to cleaning epochs, g assign dummy demands to physical dummy wafers, and y order operations that share a resource. Each operation o has start s_o , completion $e_o = s_o + d_o$, and duration d_o . The constants $\delta_{oo'}$ and $\rho_{oo'}^r$ are technological and endpoint-dependent setup separations, and M is a

valid scheduling-horizon bound.

The central model relations are compactly represented by

$$\sum_{k \in \mathcal{K}(p)} x_{pk} = 1 \quad p \in \mathcal{P}, \quad (1)$$

$$x_{pk} \leq u_k, \quad u_{k+1} \leq u_k, \quad k \in \mathcal{K}, \quad (2)$$

$$s_{o'} \geq e_o + \delta_{oo'}, \quad (o, o') \in \mathcal{A}, \quad (3)$$

$$s_{o'} \geq e_o + \rho_{oo'}^r - M(1 - y_{oo'}^r) - M(2 - b_o - b_{o'}), \quad \{o, o'\} \subseteq \mathcal{O}_r, \quad (4)$$

$$s_o \geq e_{o'} + \rho_{o'o}^r - M y_{o'o}^r - M(2 - b_o - b_{o'}), \quad \{o, o'\} \subseteq \mathcal{O}_r, \quad (5)$$

$$\sum_{\tau \in \mathcal{T}_c} g_{j\tau} = b_j \quad j \in \mathcal{J}_c^{\text{DW}}. \quad (6)$$

Equation (1) makes every product group part of exactly one recipe position. Equation (2) links assignments to a contiguous active structure. Technological arcs in Eq. (3) create complete product and dummy routes. The paired activation-qualified disjunctions (4)–(5) prevent overlap on the ATR, aligner, load-lock slots, VTR, chambers, and each reusable dummy wafer; the setup term ρ includes endpoint-dependent robot repositioning.

Table 2 connects the remaining industrial rules to their mathematical role. This mapping is important because feasibility is a property of the complete equipment trajectory, not only of chamber occupancy.

Cleaning epochs satisfy

$$\sum_e a_{vce} = b_v, \quad a_{vce} \leq z_{ce}, \quad \sum_v a_{vce} \leq H z_{ce}, \quad (7)$$

where H is the maximum number of process cycles between cleans and z_{ce} activates epoch e . Linking constraints insert a dummy-only clean between consecutive epochs and allow a 2×2 chain to split only at an exposure boundary. Residence limits take the form

$$0 \leq s_{\text{out}(w,r)} - e_{\text{in}(w,r)} \leq \bar{R}_{wr}. \quad (8)$$

In this setting, a scheduling deadlock is a circular wait in which active wafer moves hold resources while requesting one another's next resource. The MILP excludes such a cycle from every generated schedule. Every active route has a finite source-to-return precedence chain, and every pair sharing a unit-capacity resource receives exactly one temporal order. A directed cycle in the combined route–resource graph would make the corresponding

Table 2: Industrial rules represented by the MILP.

Industrial rule	Mathematical mechanism	Scheduling consequence
Rotary side sequence	Event equalities and precedence arcs	Enforces load-rotate-process-unload physics
Mixed $4 \times 1/2 \times 2$ recipes	Assignment, activation, and adjacency links	Selects feasible grouping and shared exposures
Cleaning interval	Cycle-to-epoch assignment and capacity	Inserts dummy-only cleaning at legal boundaries
Reusable dummy inventory	Physical-wafer assignment plus interval disjunction	Prevents simultaneous reuse of one dummy wafer
Robot/load-lock capacity	Endpoint-dependent resource ordering	Prevents transfer conflicts and includes repositioning
Residence restrictions	Upper bounds between paired route events	Protects time-sensitive products and transfers
Finite-batch throughput	Latest return-to-load-port variable	Measures end-to-end rather than chamber-only completion

positive separation inequalities sum to $0 > 0$ and is therefore infeasible. Terminal return constraints additionally require every product and assigned dummy wafer to complete its route. Thus each feasible solution is a complete,
 170 deadlock-free offline equipment trajectory under the modeled conditions.

3.3. Throughput first, schedule quality second

Stage one minimizes the return time of the last product wafer:

$$C_{\max} \geq e_w^{\text{return}} \quad (w \in \mathcal{W}), \quad \min_{z \in \mathcal{F}} C_{\max}(z), \quad (9)$$

where $\mathcal{F}$ contains all material, rotary, cleaning, resource, and timing constraints. A makespan objective alone is insufficient operationally. Once the
 175 last return time is fixed, many schedules may be equivalent in $C_{\max}$ but differ substantially in intermediate product waiting, extreme holds, and chamber cadence. Those differences affect exposure consistency, work-in-process residence, and recovery from small disturbances.

Let $\hat{z}$ be the best stage-one schedule and $\hat{C}$ its makespan. Stage two
 180 retains its discrete material and resource structure $\mathcal{D}_{\text{fix}}$ and imposes

$$z_d = \hat{z}_d \quad (d \in \mathcal{D}_{\text{fix}}), \quad C_{\max}(z) \leq \hat{C} + \varepsilon_C. \quad (10)$$

It then minimizes

$$\Phi(z) = \lambda_1 L_{\text{sum}} + \lambda_2 L_{\text{max}} + \lambda_3 D_{\text{cadence}} + \lambda_4 E_{\text{cap}}, \quad (11)$$

where $\mathcal{L}(z)$ contains the model-native wait and idle terms that are tight under the linear objective, and $h_{mce\ell}$ is the ℓ th active processing start for recipe mode m , chamber c , and cleaning epoch e :

$$L_{\text{sum}} = \sum_{\ell \in \mathcal{L}(z)} \ell, \quad L_{\text{max}} = \max_{\ell \in \mathcal{L}(z)} \ell, \quad (12)$$

$$D_{\text{cadence}} = \sum_{m,c,e,\ell \geq 2} |(h_{mce,\ell+1} - h_{mce,\ell}) - (h_{mce,\ell} - h_{mce,\ell-1})|, \quad (13)$$

$$E_{\text{cap}} = \sum_{\ell \in \mathcal{L}(z)} [\ell - \bar{\ell}]_+. \quad (14)$$

185 The nonnegative weights $\lambda_1, \dots, \lambda_4$ set the trade-off among aggregate waiting, extreme waiting, cadence irregularity, and soft wait-cap excess. The formal

linear profile uses weights $(1, 10, 25, 0)$; the last term is available only when a positive soft cap is configured. Physical product-route waits reported as operational outcomes are reconstructed uniformly from event times, rather than read from these objective auxiliaries; their six-segment definition is given in the supplement.

Here ε_C is only the numerical equality tolerance. Equation (10) gives the two-stage model a clear guarantee. The second stage cannot worsen the retained makespan beyond this tolerance and cannot change the chosen recipe, assignment, cleaning, or resource-order structure. If stage one proves optimality, the refined schedule remains globally makespan-optimal within the same tolerance; its timing is optimal for Eq. (11) within the retained discrete structure. If stage one ends with an incumbent, the same statement holds relative to that incumbent. The model therefore separates objectives without disconnecting them from the same physical schedule.

3.4. Modeling advance

The formulation extends existing rotating or multi-space models in four directions simultaneously: finite rather than steady-state scheduling; two material sources rather than one; endogenous mixed-recipe grouping rather than a fixed sequence; and conditional schedule-quality refinement rather than makespan alone. This combination matters because dummy availability, cleaning, and rotary-side decisions alter the robot conflict graph itself. They cannot be appended as independent capacity constraints after a product sequence has been selected.

The same structure explains computational difficulty. Assignment, activation, epoch, dummy identity, and pairwise order binaries interact with a dense event time network. The resulting cuts may affect different physical roles even when their generic efficacy scores are similar. This observation motivates the solver design in the next section.

4. Structure-aware learning-enhanced exact solution

4.1. From model structure to solver guidance

The MILP exhibits two practical symptoms: finding the first competitive incumbent can be slow, and generic cut scores do not reveal whether a cut tightens discrete equipment choices, continuous timing, or their coupling.

220 RDMCT-HBS addresses the symptoms at complementary points. Structure-preserving binary-skeleton (SPBS) initialization supplies primal guidance before branch-and-cut, while the learned selector supplies dual guidance during separation. Both components remain advisory: SCIP keeps the original feasible region, validates incumbents, generates valid cuts, and retains proof
225 responsibility.

4.2. Structure-preserving binary-skeleton warm start

SPBS fixes deterministic assignment, activation, cleaning-epoch, load-lock, and selected resource-order roles in a temporary feasibility MIP. If a profile finds no incumbent, VTR, load-lock, and other order fixings are
230 progressively released; continuous event times always remain free. A start is accepted only after reconstruction and complete feasibility checking in the full model. Only that feasible schedule—not the temporary fixing constraints—is injected into the unrestricted production model.

This design is stronger than an arbitrary primal hint. The retained decisions encode the combinatorial architecture that the detailed timing model
235 must otherwise discover from scratch, while the completion solve repairs local timing interactions. Because the original model receives only a feasible incumbent, SPBS cannot remove a better solution or weaken an optimality proof.

240 4.3. Variable-role reconstruction and hierarchical decoding

HEM describes candidate cuts with 13 generic statistics. We add ten structural features derived from solver metadata. Variables are partitioned into binary, general integer, objective-active continuous, auxiliary continuous, and other roles. For cut i , the coefficient mass on role g is

$$p_{ig} = \frac{\sum_{j:x_j \in g} |a_{ij}|}{\sum_j |a_{ij}| + \varepsilon}. \quad (15)$$

245 Role masses, bounded-variable mass, coefficient sign balance, discrete-continuous coupling, dominant-role mass, and role entropy extend the representation to 23 dimensions. The features are invariant to variable names but distinguish the physical function of a cut.

A high-level policy selects the cut budget and a low-level pointer policy
250 orders candidates (Vinyals et al., 2015). If k cuts are selected from state S ,

their joint probability is

$$\pi_{\theta}(k, i_{1:k} \mid S) = \pi_{\theta}^H(k \mid S) \prod_{t=1}^k \pi_{\theta}^L(i_t \mid S, i_{1:t-1}). \quad (16)$$

Two workers collect independent SCIP trajectories in parallel; their batches then drive a central update of the shared policy with return $R = -\text{PDI}$. The estimator

$$\widehat{\nabla_{\theta} J} = \frac{1}{B} \sum_{b=1}^B (R_b - \bar{R}_{\text{EMA}}) \nabla_{\theta} \log \pi_{\theta}(\tau_b), \quad \bar{R}_{\text{EMA}} \leftarrow 0.9 \bar{R}_{\text{EMA}} + 0.1 \bar{R} \quad (17)$$

255 uses an exponential-moving-average return baseline, not an independently learned critic. We therefore describe the backbone as A3C-inspired parallel policy gradient (Mnih et al., 2016); the new contributions are the scheduling warm start, role reconstruction, decoding, and completion mechanisms.

At inference, bounded beam search keeps alternative partial sequences
 260 instead of committing to the first locally best cut (Choo et al., 2022). Role-submodular completion then adds candidates that balance policy score, immediate cut quality, role coverage, and representation of the remaining pool (Nemhauser et al., 1978). The mechanisms are complementary: reconstruction supplies meaningful diversity; beam decoding explores alternative orders;
 265 completion prevents the final set from collapsing onto one variable role.

4.4. *Exactness and computational scope*

RDMCT-HBS ranks only cuts already generated and certified as valid by SCIP (Achterberg, 2009; Bestuzheva et al., 2023). SPBS adds only a solver-checked incumbent. Consequently, neither component changes the objective,
 270 constraints, integrality requirements, or lower-bound logic. The method can improve anytime behavior while preserving the same mathematical meaning of feasibility and optimality as default SCIP. The supplementary material documents the reproducibility protocol.

5. Computational results and analysis

275 5.1. *Evidence structure*

The study uses seven complementary evidence sets. Four exact cases verify the physical formulation. A 2700-run nominal matrix compares seven

methods over 20 manufacturing instances and independent solver and training realizations. A 960-run matched experiment isolates SPBS. A 1500-run control isolates the second-stage mechanism and its objective. A 300-run factorial design explains how wafer count, recipe composition, and process scale shape difficulty. An 80-run one-factor study tests operational solvability under parameter perturbations. Finally, 729 runs on nine shifted 40–64-wafer configurations test larger-scale feasible-schedule delivery. All saved incumbents supporting reported conclusions pass independent checks of the corresponding physical MIP.

PDI is the primary anytime metric because it integrates SCIP’s normalized primal–dual gap over the common horizon; lower is better. Terminal gap is the normalized bound gap at termination, and reported solution time covers the complete primary solve and conditional refinement within the common wall-clock budget. Repetitions are aggregated within manufacturing instances before paired bootstrap intervals, Wilcoxon tests, and Holm adjustment are calculated. This prevents repeated random seeds from being treated as independent manufacturing evidence (Gleixner et al., 2021). Hardware, budgets, hyperparameters, complete matrix definitions, and validation tolerances are reported in the supplement.

5.2. The formulation represents the intended equipment

Table 3 reports four exact pure and mixed-flow cases. Both stages reach zero gap, and independent reconstruction confirms every route, resource, cleaning, residence, and dummy-allocation constraint.

Table 3: Exact verification of the two-stage formulation.

Configuration	Products	$C_{\max}$	Route-total wait	Route-maximum wait	Time (s)
Pure 4×1	4	458	380	76	0.009
Pure 2×2	4	584	112	28	0.039
Mixed flow	8	834	506	76	0.458
Mixed flow	16	1170	1550	99	24.390

The pure cases recover the distinct rotary logic of each recipe, while mixed cases require both recipe structures, two load-lock paths, shared robots, and dummy reuse to coexist in one feasible trajectory. The 16-product case also exercises the coupled assignment and resource-order constraints targeted by RDMCT-HBS.

5.3. RDMCT-HBS attains the lowest observed nominal mean PDI and time

Table 4 gives the full nominal comparison. RDMCT-HBS has the lowest mean PDI and the shortest mean solution time among all seven methods. Its mean PDI is 1.41% below adapted HEM and 1.97% below default SCIP. Every one of its 500 runs returns a feasible incumbent.

All seven conditions receive the same independently verified automatic SPBS incumbent, time limit, memory limit, and solver seeds. The comparison therefore isolates cut-selection and decoding choices rather than the total SPBS effect, which is identified separately below. SCIP denotes its default cut selector under this shared start; adaptive cut selection (ACS) uses a tuned four-score SCIP rule. HEM-13D and HEM-13D+Beam share the original 13-feature policy and differ only in decoding. The 23D-only condition uses reconstructed features without structural reranking; Structure+Greedy adds role-based completion; and RDMCT-HBS combines that structure with beam decoding.

Table 4: Nominal comparison under a common SPBS start; statistics are first averaged within 20 manufacturing instances.

Method	Mean PDI	Mean gap (%)	Mean time (s)	Optimal (%)
SCIP default selector	16919.01	119.54	226.05	64.0
Adaptive cut selection (ACS)	16854.87	127.18	222.34	62.0
HEM-13D	16823.18	127.24	219.65	62.4
HEM-13D+Beam	16829.59	127.19	219.95	62.4
23D Features Only	16620.91	125.33	219.42	62.6
Structure+Greedy	16598.11	126.03	218.44	63.2
RDMCT-HBS	16585.70	126.01	218.32	63.2

The pattern across methods explains the advantage. Replacing HEM’s generic representation with variable-role features produces the largest component gain. Adding structural completion improves the mean again, and beam decoding provides a further improvement only after the representation and completion mechanism are structure-aware. Beam alone slightly worsens HEM-13D. The result shows that the benefit does not come from wider decoding in isolation; it comes from exploring alternatives in a representation that distinguishes the roles those alternatives play in the scheduling MILP.

SCIP has a marginally better final gap and optimal rate, whereas RDMCT-HBS leads PDI and elapsed time. This is an interpretable trade-off rather than a contradiction. SCIP’s default selection remains

strong for eventual proof, while RDMCT-HBS improves the trajectory by allocating cuts toward the coupled assignment–timing structure earlier. For time-constrained industrial use, where an incumbent and its quality
335 during the available horizon matter, RDMCT-HBS has the lowest observed nominal mean PDI and time. The instance-first mean PDI improvements over HEM and SCIP are 237.49 (95% bootstrap CI $[-446.71, 846.26]$) and 333.31 (CI $[-382.16, 999.70]$), respectively; both Holm-adjusted tests have $p = 1$. Thus the nominal ranking is an observed average advantage over this
340 matrix, whereas the statistically supported controlled gains are the SPBS and second-stage effects reported next.

5.4. SPBS provides a significant primal improvement

The independent SPBS experiment gives the clearest mechanism-level result. As Table 5 shows, automatic SPBS raises incumbent availability
345 from 51.67% to 100% and reduces mean PDI by 26.19%. The effects remain significant after Holm adjustment ($p = 2.55 \times 10^{-5}$ and $p = 3.46 \times 10^{-8}$, respectively).

Table 5: Independent SPBS comparison over 48 instances and ten matched seeds.

Warm start	Runs	Incumbent (%)	Optimal (%)	Mean PDI
Automatic SPBS	480	100.00	54.58	20278.78
None	480	51.67	51.67	27473.05

The industrial interpretation is direct. Without a structured start, nearly half of the runs end before a complete schedule is available. SPBS converts
350 the equipment’s discrete architecture into an incumbent early enough for SCIP to improve it and use it to prune the tree. The large PDI reduction therefore reflects both better primal availability and a stronger interaction between the incumbent and subsequent exact search. SPBS is not merely a convenience for reporting a feasible schedule; it changes the quality of the
355 entire anytime trajectory.

5.5. The second stage creates operationally superior schedules

Makespan verification alone cannot establish the value of the second stage. We therefore reconstruct route waits directly from product event times and cadence directly from active chamber starts. This avoids using objective slack

Table 6: Operational effect of conditional schedule refinement.

Profile	Route-total wait	Route-maximum wait	Cadence CV	Cadence deviation
Full two-stage model	2596.01	296.33	0.130	22.04
Quadratic objective (linear off)	2469.47	293.27	0.134	60.18
No second stage	5398.60	952.57	0.216	370.41

variables as performance measures and makes the metric identical across all profiles.

Relative to no second stage, the proposed refinement reduces route-total wait by 51.91%, maximum route wait by 68.89%, cadence CV by 39.70%, and cadence deviation by 94.05%; all four effects remain significant after Holm adjustment. These are large operational changes, not small numerical preferences between equivalent schedules. Shorter and less extreme waits reduce residence exposure, while regular chamber starts produce a more stable rhythm for downstream handling.

The comparison with quadratic refinement adds a second insight. The quadratic profile has slightly lower route-wait means, but the proposed linear objective reduces cadence deviation by 63.38% with a significant matched effect. Thus, the proposed stage establishes the strongest cadence regularity while retaining route waits far below the no-refinement schedule. This result confirms its intended industrial role: remove severe waiting and create a repeatable processing rhythm without sacrificing throughput.

5.6. Mixed-flow scale, not uniform processing inflation, drives difficulty

The factorial results reveal a sharp complexity transition. All 8- and 16-wafer runs prove optimality, but the rate falls to 36% at 24 wafers and 20% at 32 wafers. Mean PDI increases from 28.13 and 1876.54 to 31155.94 and 32959.89, respectively.

Table 7: Computational transition by wafer count.

Wafers	Runs	Optimal (%)	Mean PDI	Mean time (s)
8	75	100	28.13	0.34
16	75	100	1876.54	33.86
24	75	36	31155.94	408.55
32	75	20	32959.89	433.20

385 The fitted factorial model ($R^2 = 0.998$ for PDI) attributes this transition to the interaction of wafer count and recipe composition. At the 1:1 mix, the additional PDI interactions are 37729.48 and 40788.63 for 24 and 32 wafers. By contrast, every prespecified processing-scale main and interaction interval includes zero. The difficult cases are therefore created by combinatorial coupling—more products participating in both recipe structures—rather than by uniformly longer processing durations. This finding supports the design of RDMCT-HBS: structural assignment and timing roles are more informative than a generic measure of problem duration.




Figure 2: Prespecified factorial effects on PDI. Points are OLS coefficients, bars are 95% confidence intervals, and the vertical line denotes zero.

390 *5.7. Sensitivity confirms operational solvability*

All 80 one-factor sensitivity runs return independently valid incumbents. The lowest mean PDI occurs with a longer cleaning interval (28899.41), whereas slower motion produces the highest mean PDI (42478.32). The pattern is physically consistent: less frequent cleaning removes dummy-only
 395 interruptions, while slower transport intensifies conflicts on already shared robots. More importantly, the formulation and exact-solution chain remain operational across every prescribed perturbation.

5.8. OOD stress test confirms scalable feasibility

The distribution-shift suite enlarges the batch to 40–64 wafers and
 400 changes the product-mixture composition. Every one of the 729 runs returns a saved incumbent that passes independent validation. Because every run uses the complete search horizon, PDI and terminal gap provide the discriminating measures in Table 8.

Table 8: OOD stress test; statistics are first averaged within nine manufacturing instances.

Method	Runs	Incumbent (%)	Mean PDI	Mean gap (%)
SCIP	27	100	73670.99	171.35
Adaptive cut selection	27	100	70915.54	160.78
HEM-13D	135	100	73003.59	165.88
HEM-13D+Beam	135	100	73165.58	167.04
23D Features Only	135	100	74187.22	169.96
Structure+Greedy	135	100	73507.76	168.54
RDMCT-HBS	135	100	73919.15	170.02

Adaptive cut selection records the lowest OOD PDI, while RDMCT-HBS
 405 has an observed mean terminal gap 0.78% below default SCIP. No OOD pairwise contrast remains significant after global Holm adjustment. The complete outcome separates two forms of evidence: RDMCT-HBS leads observed mean anytime search in the nominal production range, and the shared SPBS–MILP–SCIP pipeline delivers a validated feasible schedule in every
 410 larger shifted run. All seven methods use the full search horizon and none proves OOD optimality, identifying global bound closure as the remaining extreme-scale challenge.

5.9. Industrial implications

Three implications follow from the complete evidence. First, throughput and schedule quality should be optimized sequentially: a makespan-only model leaves substantial avoidable waiting and cadence instability. Second, an application-feasible incumbent gives the largest controlled computational improvement measured here; SPBS should therefore be treated as part of the solution architecture rather than a preprocessing detail. Third, learning is most effective when its representation mirrors the MIP’s physical roles. The observed sign reversal of the beam effect is consistent with coupling between beam search and the structure-aware representation: beam decoding improves the observed mean only after candidate roles are aligned with discrete equipment choices and continuous timing. This comparison does not itself test a statistical interaction.

The explicit state, resource, and timing layers also provide a direct foundation for online arrivals, stochastic processing, and module-failure recourse. The present finite-batch evidence establishes the central result: for dual-source mixed-flow equipment, the two-stage model produces materially better schedules and RDMCT-HBS attains the lowest observed nominal mean PDI and time among the compared methods.

6. Conclusions

This paper develops an integrated model-and-algorithm framework for dual-source mixed-flow rotating-chamber cluster tools. The two-stage MILP captures the complete interaction of product wafers, reusable vacuum-zone dummy wafers, mixed rotary recipes, cleaning, shared robots, load locks, and residence restrictions. Its first stage optimizes end-to-end makespan; its second stage preserves that production plan and transforms its timing quality.

The experimental results establish the value of both contributions. The full model yields physically valid pure and mixed schedules, and second-stage refinement cuts route-total wait, maximum route wait, cadence variation, and cadence deviation by 51.91%, 68.89%, 39.70%, and 94.05%. The proposed RDMCT-HBS solver records the lowest observed mean PDI and solution time in the complete 2700-run seven-method comparison. Its SPBS component independently reduces PDI by 26.19% and raises incumbent availability to 100%. Factorial results further show that the principal difficulty is the interaction between scale and mixed-recipe composition, the same class of

discrete–continuous coupling represented by the reconstructed features. All
450 729 OOD runs also produce independently validated incumbents, establishing
feasible schedule delivery on nine larger shifted configurations.

The central conclusion is therefore unambiguous: the proposed two-stage
MILP produces operationally superior schedules, and the structure-aware
learning-enhanced branch-and-cut method attains the best observed nomi-
455 nal mean anytime performance in this comparison. Together they provide
a rigorous and practically useful approach to a class of rotating-chamber
scheduling problems that is not covered by existing single-source or fixed-
sequence models.

Declaration of competing interest

460 The authors declare that they have no known competing financial inter-
ests or personal relationships that could have appeared to influence the work
reported in this paper.

Data and code availability

The data, mathematical instances, trained models, and source code sup-
465 porting the findings are available from the corresponding author on reason-
able request. A permanent public archive will be cited in the version of
record.

Declaration of generative AI and AI-assisted technologies in the writing process

470 During the preparation of this work, the authors used OpenAI Codex
for language editing, consistency checking, and code review. The authors
reviewed and edited all AI-assisted output and take full responsibility for the
content of the publication.

References

475 Achterberg, T., 2009. SCIP: Solving constraint integer programs.
Mathematical Programming Computation 1, 1–41. doi:10.1007/
s12532-008-0001-1.

- Bengio, Y., Lodi, A., Prouvost, A., 2021. Machine learning for combinatorial optimization: A methodological tour d’horizon. *European Journal of Operational Research* 290, 405–421. doi:10.1016/j.ejor.2020.07.063.
- Bestuzheva, K., Besançon, M., Chen, W.K., Chmiela, A., Donkiewicz, T., van Doornmalen, J., Eifler, L., Gaul, O., Gamrath, G., Gleixner, A., Gottwald, L., Graczyk, C., Halbig, K., Hoen, A., Hojny, C., van der Hulst, R., Koch, T., Lübbecke, M.E., Maher, S.J., Matter, F., Mühmer, E., Müller, B., Pfetsch, M.E., Schlösser, F., Serrano, F., Shinano, Y., Sofranac, B., Turner, M., Vigerske, S., Wegscheider, F., Wellner, P., Weninger, D., Witzig, J., 2023. Enabling research through the SCIP optimization suite 8.0. *ACM Transactions on Mathematical Software* 49, 22:1–22:21. doi:10.1145/3585516.
- Chien, C.F., Lan, Y.B., 2021. Agent-based approach integrating deep reinforcement learning and hybrid genetic algorithm for dynamic scheduling for Industry 3.5 smart production. *Computers & Industrial Engineering* 162, 107782. doi:10.1016/j.cie.2021.107782.
- Choo, J., Kwon, Y.D., Kim, J., Jae, J., Hottung, A., Tierney, K., Gwon, Y., 2022. Simulation-guided beam search for neural combinatorial optimization, in: *Advances in Neural Information Processing Systems*. URL: https://papers.nips.cc/paper_files/paper/2022/hash/39b9b60f0d149eabd1fff2d7c7d5afc4-Abstract-Conference.html, doi:10.52202/068431-0637.
- Gleixner, A., Hendel, G., Gamrath, G., Achterberg, T., Bastubbe, M., Berthold, T., Christophel, P., Jarck, K., Koch, T., Linderoth, J., Lübbecke, M.E., Mittelman, H.D., Ozyurt, D., Ralphs, T.K., Salvagnin, D., Shinano, Y., 2021. MIPLIB 2017: Data-driven compilation of the 6th mixed-integer programming library. *Mathematical Programming Computation* 13, 443–490. doi:10.1007/s12532-020-00194-3.
- Gu, L., Wu, N., Qiao, Y., Zhang, S., 2024. Scheduling of Multi-finger-robotic Cluster Tools with Multi-space Process Modules, in: *2024 International Conference on Networking, Sensing and Control (ICNSC)*, IEEE. doi:10.1109/ICNSC62968.2024.10760109.
- Hong, T.Y., Hsu, C.Y., Chen, C.L., 2026. Simulation-based deep reinforcement learning for process time optimization in semiconductor cluster tool

and an empirical study. *Computers & Industrial Engineering* 211, 111698. doi:10.1016/j.cie.2025.111698.

515 Huang, Z., Wang, K., Liu, F., Zhen, H.L., Zhang, W., Yuan, M., Hao, J., Yu, Y., Wang, J., 2022. Learning to select cuts for efficient mixed-integer programming. *Pattern Recognition* 123, 108353. doi:10.1016/j.patcog.2021.108353.

Khalil, E.B., Morris, C., Lodi, A., 2022. MIP-GNN: A data-driven framework for guiding combinatorial solvers, in: *Proceedings of the AAAI Conference on Artificial Intelligence*, pp. 10219–10227. doi:10.1609/aaai.v36i9.21262.

Kim, D., Kim, H.J., 2025. Noncyclic scheduling of cluster tools using deep reinforcement learning. *IEEE Transactions on Automation Science and Engineering* 22, 24009–24023. doi:10.1109/TASE.2025.3632534.

525 Kim, H.J., Lee, J.H., 2025. Scheduling cluster tools for concurrent processing: Deep reinforcement learning with adaptive search. *IEEE Transactions on Automation Science and Engineering* 22, 3783–3796. doi:10.1109/TASE.2024.3399818.

530 Kim, W., Yu, T.S., Lee, T.E., 2021. Integrated scheduling of a dual-armed cluster tool for maximizing steady schedule patterns. *IEEE Transactions on Systems, Man, and Cybernetics: Systems* 51, 7282–7294. doi:10.1109/TSMC.2020.2978486.

Ko, S.G., Yu, T.S., Lee, T.E., 2021. Wafer delay analysis and workload balancing of parallel chambers for dual-armed cluster tools with multiple wafer types. *IEEE Transactions on Automation Science and Engineering* 18, 1516–1526. doi:10.1109/TASE.2021.3061140.

Lee, T.G., Yu, T.S., Lee, T.E., 2023. Cleaning plan optimization for dual-armed cluster tools with general chamber cleaning periods. *IEEE Transactions on Automation Science and Engineering* 20, 1890–1906. doi:10.1109/TASE.2022.3188858.

540 Li, C., Chang, Q., 2022. Hybrid feedback and reinforcement learning-based control of machine cycle time for a multi-stage production system. *Journal of Manufacturing Systems* 65, 351–361. doi:10.1016/j.jmsy.2022.09.020.

- 545 Li, C., Yang, F., Zhen, L., 2022. Efficient scheduling approaches to time-constrained single-armed cluster tools with condition-based chamber cleaning operations. *International Journal of Production Research* 60, 3555–3568. doi:10.1080/00207543.2021.1926568.
- Li, X., Yang, W., 2025. Cyclic scheduling of multi-type wafers concurrent
550 processing in single-arm cluster tools with residency time constraints. *Expert Systems with Applications* 269, 126443. doi:10.1016/j.eswa.2025.126443.
- Liu, J., Qiao, F., Zou, M., Zinn, J., Ma, Y., Vogel-Heuser, B., 2022. Dynamic scheduling for semiconductor manufacturing systems with uncertainties
555 using convolutional neural networks and reinforcement learning. *Complex & Intelligent Systems* 8, 4641–4662. doi:10.1007/s40747-022-00844-0.
- Mnih, V., Badia, A.P., Mirza, M., Graves, A., Lillicrap, T.P., Harley, T., Silver, D., Kavukcuoglu, K., 2016. Asynchronous methods for deep reinforcement learning, in: *Proceedings of the 33rd International Conference on Machine Learning*, PMLR. pp. 1928–1937. URL: <https://proceedings.mlr.press/v48/mniha16.html>.
560
- Nemhauser, G.L., Wolsey, L.A., Fisher, M.L., 1978. An analysis of approximations for maximizing submodular set functions—i. *Mathematical Programming* 14, 265–294. doi:10.1007/BF01588971.
- 565 Pan, C., Lai, Y., Qiao, Y., Wu, N., Luo, X., 2024. Cycle time analysis for cluster tools with parallel process chambers, processing time variation, and chamber cleaning operations. *Expert Systems with Applications* 255, 124471. doi:10.1016/j.eswa.2024.124471.
- Paulus, M.B., Zarpellon, G., Krause, A., Charlin, L., Maddison, C., 2022. Learning to cut by looking ahead: Cutting plane selection via imitation
570 learning, in: *Proceedings of the 39th International Conference on Machine Learning*, PMLR. pp. 17584–17600. URL: <https://proceedings.mlr.press/v162/paulus22a.html>.
- Qiao, Y., Lu, Y., Li, J., Zhang, S., Wu, N., Liu, B., 2022. An efficient binary integer programming model for residency time-constrained cluster
575 tools with chamber cleaning requirements. *IEEE Transactions on Automation Science and Engineering* 19, 1757–1771. doi:10.1109/TASE.2021.3122576.

- Scavuzzo, L., Aardal, K., Lodi, A., Yorke-Smith, N., 2024. Machine learning augmented branch and bound for mixed integer linear programming. Mathematical Programming doi:10.1007/s10107-024-02130-y.
- 580
- Tang, Y., Agrawal, S., Faenza, Y., 2020. Reinforcement learning for integer programming: Learning to cut, in: Proceedings of the 37th International Conference on Machine Learning, PMLR. pp. 9367–9376. URL: <https://proceedings.mlr.press/v119/tang20a.html>.
- 585
- Turner, M., Koch, T., Serrano, F., Winkler, M., 2023. Adaptive cut selection in mixed-integer linear programming. Open Journal of Mathematical Optimization 4, 1–28. doi:10.5802/ojmo.25.
- Vinyals, O., Fortunato, M., Jaitly, N., 2015. Pointer networks, in: Advances in Neural Information Processing Systems, pp. 2692–2700.
- 590
- Wang, J., Liu, C., Zhou, M., Leng, T., Albeshri, A., 2023a. Optimal cyclic scheduling of wafer-residency-time-constrained dual-arm cluster tools by configuring processing modules and robot waiting time. IEEE Transactions on Semiconductor Manufacturing 36, 251–259. doi:10.1109/TSM.2023.3239198.
- 595
- Wang, J., Wang, Z., Li, X., Kuang, Y., Shi, Z., Zhu, F., Yuan, M., Zeng, J., Zhang, Y., Wu, F., 2024. Learning to cut via hierarchical sequence/set model for efficient mixed-integer programming. IEEE Transactions on Pattern Analysis and Machine Intelligence 46, 9697–9713. doi:10.1109/TPAMI.2024.3432716.
- 600
- Wang, Z., Li, X., Wang, J., Kuang, Y., Yuan, M., Zeng, J., Zhang, Y., Wu, F., 2023b. Learning cut selection for mixed-integer linear programming via hierarchical sequence model, in: Proceedings of the Eleventh International Conference on Learning Representations. URL: <https://openreview.net/forum?id=Zob4P9bRNcK>.
- 605
- Yang, F., Li, C., Zhen, L., Zhang, C., Wu, N., 2025. Scheduling residency time-constrained single-armed cluster tools with two-wafer types via mixed integer programming model. Expert Systems with Applications 288, 128209. doi:10.1016/j.eswa.2025.128209.

- 610 Yuan, F., Zhao, Q., Huang, B., Pan, C., 2023. Scheduling of time-constrained single-arm cluster tools with purge operations in wafer fabrications. *Journal of Systems Architecture* 134, 102788. doi:10.1016/j.sysarc.2022.102788.
- Zhang, L., Lin, Y., Xu, C., Liu, M., 2024. A new EDA algorithm combined
615 with Q-learning for semiconductor final testing scheduling problem. *Computers & Industrial Engineering* 193, 110259. doi:10.1016/j.cie.2024.110259.
- Zhu, Q., Li, B., Hou, Y., Li, H., Wu, N., 2023. Scheduling dual-arm multi-cluster tools with regulation of post-processing time. *IEEE/CAA Journal of Automatica Sinica* 10, 1730–1742. doi:10.1109/JAS.2023.123189.
620
- Zhu, Q., Wang, G., Hou, Y., Wu, N., Qiao, Y., 2022. Optimally scheduling dual-arm multi-cluster tools to process two wafer types. *IEEE Robotics and Automation Letters* 7, 5920–5927. doi:10.1109/LRA.2022.3157031.